

# Integrated Project 2

## *Recovered gold share prediction model*

### Table of Contents

- [1 Goal](#)
- [2 Data description](#)
- [3 Feature naming](#)
- [4 Imports](#)
- [5 Input data](#)
- [6 Descriptive statistics](#)
- [7 Recovery calculation check](#)
- [8 Missing features analysis](#)
- [9 Preprocessing](#)
  - [9.1 Data type change](#)
  - [9.2 Test targets](#)
  - [9.3 Missing values](#)
  - [9.4 Duplicates](#)
- [10 EDA](#)
  - [10.1 Concentrations of metals](#)
  - [10.2 Feed particle size comparison](#)
  - [10.3 Outliers](#)
  - [10.4 Target analysis](#)
- [11 Common columns](#)
- [12 Standard scaling](#)
- [13 Evaluation metric](#)
- [14 Model selection](#)
- [15 Hyperparameter tuning](#)
- [16 Test the model](#)
- [17 Feature importance](#)
- [18 Sanity check](#)

### Goal

Develop a machine learning model for Zyfra, efficiency solutions for heavy industry company, that would analyze data on extraction and purification of gold from an ore and predict the amount of recovered gold. The model will help to optimize the production and eliminate unprofitable parameters.

# Data description

## Technological process

- *Rougher feed* — raw material
- *Rougher additions* (or *reagent additions*) — flotation reagents: *Xanthate*, *Sulphate*, *Depressant*
  - *Xanthate* — promoter or flotation activator;
  - *Sulphate* — sodium sulphide for this particular process;
  - *Depressant* — sodium silicate.
- *Rougher process* — flotation
- *Rougher tails* — product residues
- *Float banks* — flotation unit
- *Cleaner process* — purification
- *Rougher Au* — rougher gold concentrate
- *Final Au* — final gold concentrate

## Parameters of stages

- *air amount* — volume of air
- *fluid levels*
- *feed size* — feed particle size
- *feed rate*

# Feature naming

Here's how we name the features:

```
[stage].[parameter_type].[parameter_name]
```

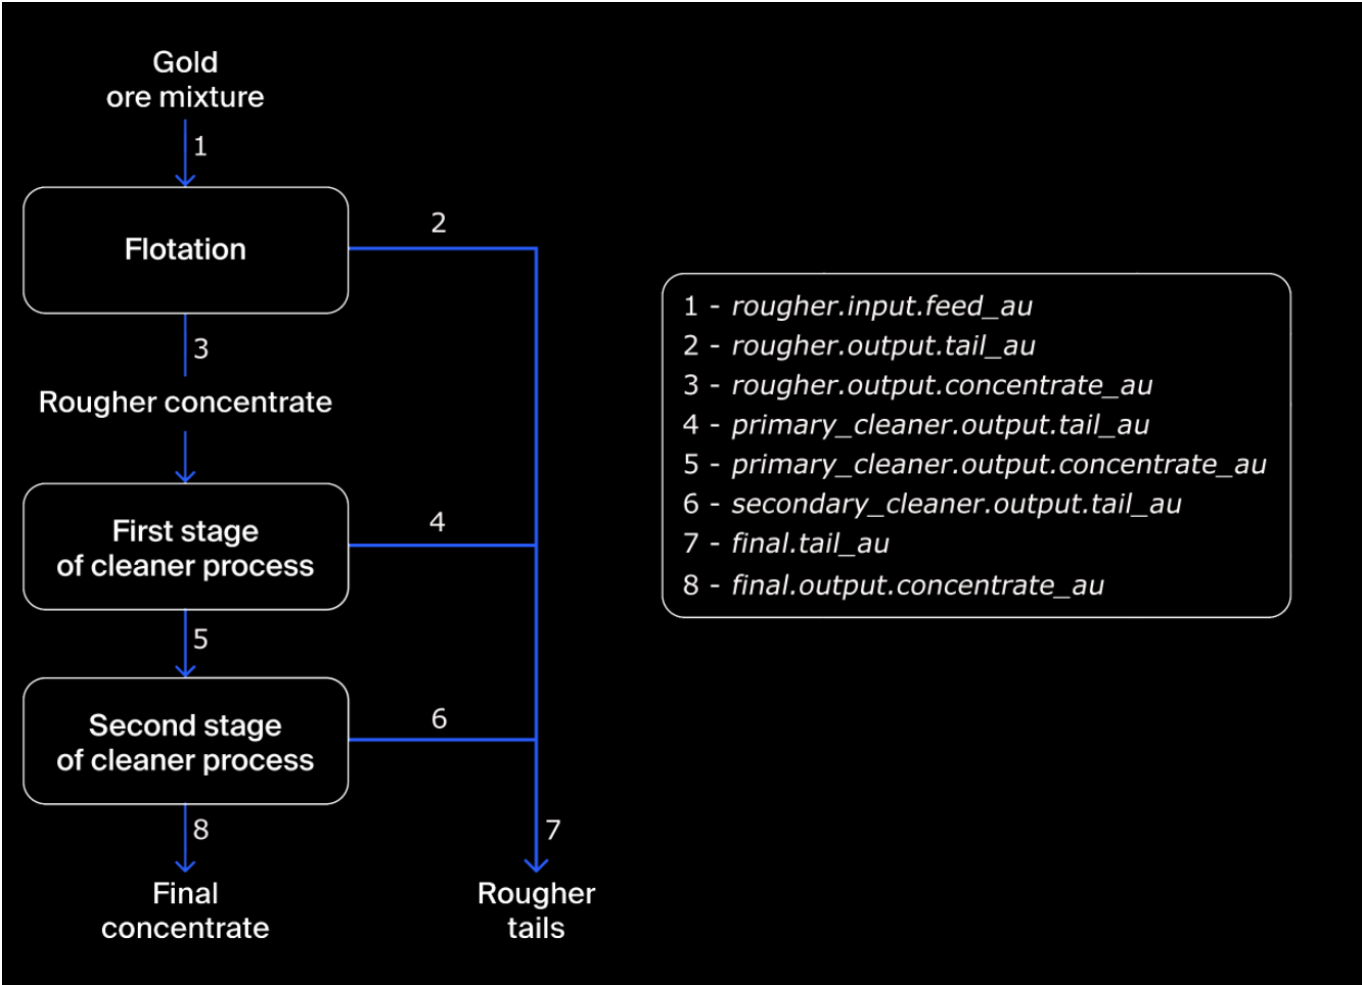
Example: `rougher.input.feed_ag`

Possible values for `[stage]` :

- *rougher* — flotation
- *primary\_cleaner* — primary purification
- *secondary\_cleaner* — secondary purification
- *final* — final characteristics

Possible values for `[parameter_type]` :

- *input* — raw material parameters
- *output* — product parameters
- *state* — parameters characterizing the current state of the stage
- *calculation* — calculation characteristics



Imports

In [2]:

```
import pandas as pd
import matplotlib
import numpy as np
import seaborn as sns

from sklearn.preprocessing import StandardScaler as ss

from sklearn.dummy import DummyRegressor
from sklearn import linear_model
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn import svm

from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer

import matplotlib.pyplot as plt
%matplotlib inline

import sys
import warnings
if not sys.warnoptions:
    warnings.simplefilter("ignore")

pd.set_option('display.max_rows', None, 'display.max_columns', None)

print("Setup Complete")
```

Setup Complete

## Input data

In [3]:

```
try:
    df_train = pd.read_csv('gold_recovery_train.csv')
    df_test = pd.read_csv('gold_recovery_test.csv')
    df_full = pd.read_csv('gold_recovery_full.csv')

except:
    df_train = pd.read_csv('/datasets/gold_recovery_train.csv')
    df_test = pd.read_csv('/datasets/gold_recovery_test.csv')
    df_full = pd.read_csv('/datasets/gold_recovery_full.csv')
```

## Descriptive statistics

In [4]:

```
df_train.head()
```

Out[4]:

	date	final.output.concentrate_ag	final.output.concentrate_pb	final.output.concentrate_sol
0	2016-01-15 00:00:00	6.055403	9.889648	5.507324
1	2016-01-15 01:00:00	6.029369	9.968944	5.257781
2	2016-01-15 02:00:00	6.055926	10.213995	5.383759
3	2016-01-15 03:00:00	6.047977	9.977019	4.858634
4	2016-01-15 04:00:00	6.148599	10.142511	4.939416

In [5]:

```
df_test.head()
```

Out[5]:

	date	primary_cleaner.input.sulfate	primary_cleaner.input.depressant	primary_cleaner.input
0	2016-09-01 00:59:59	210.800909	14.993118	
1	2016-09-01 01:59:59	215.392455	14.987471	
2	2016-09-01 02:59:59	215.259946	12.884934	
3	2016-09-01 03:59:59	215.336236	12.006805	
4	2016-09-01 04:59:59	199.099327	10.682530	

Notes for preprocessing:

- All features are numerical, except for the `date` variable.

In [6]:

```
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 16860 entries, 0 to 16859
```

```
Data columns (total 87 columns):
```

#	Column	Non-Null Count	Dtype
0	date	16860 non-null	object
1	final.output.concentrate_ag	16788 non-null	float64
2	final.output.concentrate_pb	16788 non-null	float64
3	final.output.concentrate_sol	16490 non-null	float64
4	final.output.concentrate_au	16789 non-null	float64
5	final.output.recovery	15339 non-null	float64
6	final.output.tail_ag	16794 non-null	float64
7	final.output.tail_pb	16677 non-null	float64
8	final.output.tail_sol	16715 non-null	float64
9	final.output.tail_au	16794 non-null	float64
10	primary_cleaner.input.sulfate	15553 non-null	float64
11	primary_cleaner.input.depressant	15598 non-null	float64
12	primary_cleaner.input.feed_size	16860 non-null	float64
13	primary_cleaner.input.xanthate	15875 non-null	float64
14	primary_cleaner.output.concentrate_ag	16778 non-null	float64
15	primary_cleaner.output.concentrate_pb	16502 non-null	float64
16	primary_cleaner.output.concentrate_sol	16224 non-null	float64
17	primary_cleaner.output.concentrate_au	16778 non-null	float64
18	primary_cleaner.output.tail_ag	16777 non-null	float64
19	primary_cleaner.output.tail_pb	16761 non-null	float64
20	primary_cleaner.output.tail_sol	16579 non-null	float64
21	primary_cleaner.output.tail_au	16777 non-null	float64
22	primary_cleaner.state.floatbank8_a_air	16820 non-null	float64
23	primary_cleaner.state.floatbank8_a_level	16827 non-null	float64
24	primary_cleaner.state.floatbank8_b_air	16820 non-null	float64
25	primary_cleaner.state.floatbank8_b_level	16833 non-null	float64
26	primary_cleaner.state.floatbank8_c_air	16822 non-null	float64

27	primary_cleaner.state.floatbank8_c_level	16833	non-n
u11	float64		
28	primary_cleaner.state.floatbank8_d_air	16821	non-n
u11	float64		
29	primary_cleaner.state.floatbank8_d_level	16833	non-n
u11	float64		
30	rougher.calculation.sulfate_to_au_concentrate	16833	non-n
u11	float64		
31	rougher.calculation.floatbank10_sulfate_to_au_feed	16833	non-n
u11	float64		
32	rougher.calculation.floatbank11_sulfate_to_au_feed	16833	non-n
u11	float64		
33	rougher.calculation.au_pb_ratio	15618	non-n
u11	float64		
34	rougher.input.feed_ag	16778	non-n
u11	float64		
35	rougher.input.feed_pb	16632	non-n
u11	float64		
36	rougher.input.feed_rate	16347	non-n
u11	float64		
37	rougher.input.feed_size	16443	non-n
u11	float64		
38	rougher.input.feed_sol	16568	non-n
u11	float64		
39	rougher.input.feed_au	16777	non-n
u11	float64		
40	rougher.input.floatbank10_sulfate	15816	non-n
u11	float64		
41	rougher.input.floatbank10_xanthate	16514	non-n
u11	float64		
42	rougher.input.floatbank11_sulfate	16237	non-n
u11	float64		
43	rougher.input.floatbank11_xanthate	14956	non-n
u11	float64		
44	rougher.output.concentrate_ag	16778	non-n
u11	float64		
45	rougher.output.concentrate_pb	16778	non-n
u11	float64		
46	rougher.output.concentrate_sol	16698	non-n
u11	float64		
47	rougher.output.concentrate_au	16778	non-n
u11	float64		
48	rougher.output.recovery	14287	non-n
u11	float64		
49	rougher.output.tail_ag	14610	non-n
u11	float64		
50	rougher.output.tail_pb	16778	non-n
u11	float64		
51	rougher.output.tail_sol	14611	non-n
u11	float64		
52	rougher.output.tail_au	14611	non-n
u11	float64		
53	rougher.state.floatbank10_a_air	16807	non-n
u11	float64		
54	rougher.state.floatbank10_a_level	16807	non-n
u11	float64		
55	rougher.state.floatbank10_b_air	16807	non-n
u11	float64		
56	rougher.state.floatbank10_b_level	16807	non-n
u11	float64		
57	rougher.state.floatbank10_c_air	16807	non-n



```

ull float64
58 rougher.state.floatbank10_c_level 16814 non-n
ull float64
59 rougher.state.floatbank10_d_air 16802 non-n
ull float64
60 rougher.state.floatbank10_d_level 16809 non-n
ull float64
61 rougher.state.floatbank10_e_air 16257 non-n
ull float64
62 rougher.state.floatbank10_e_level 16809 non-n
ull float64
63 rougher.state.floatbank10_f_air 16802 non-n
ull float64
64 rougher.state.floatbank10_f_level 16802 non-n
ull float64
65 secondary_cleaner.output.tail_ag 16776 non-n
ull float64
66 secondary_cleaner.output.tail_pb 16764 non-n
ull float64
67 secondary_cleaner.output.tail_sol 14874 non-n
ull float64
68 secondary_cleaner.output.tail_au 16778 non-n
ull float64
69 secondary_cleaner.state.floatbank2_a_air 16497 non-n
ull float64
70 secondary_cleaner.state.floatbank2_a_level 16751 non-n
ull float64
71 secondary_cleaner.state.floatbank2_b_air 16705 non-n
ull float64
72 secondary_cleaner.state.floatbank2_b_level 16748 non-n
ull float64
73 secondary_cleaner.state.floatbank3_a_air 16763 non-n
ull float64
74 secondary_cleaner.state.floatbank3_a_level 16747 non-n
ull float64
75 secondary_cleaner.state.floatbank3_b_air 16752 non-n
ull float64
76 secondary_cleaner.state.floatbank3_b_level 16750 non-n
ull float64
77 secondary_cleaner.state.floatbank4_a_air 16731 non-n
ull float64
78 secondary_cleaner.state.floatbank4_a_level 16747 non-n
ull float64
79 secondary_cleaner.state.floatbank4_b_air 16768 non-n
ull float64
80 secondary_cleaner.state.floatbank4_b_level 16767 non-n
ull float64
81 secondary_cleaner.state.floatbank5_a_air 16775 non-n
ull float64
82 secondary_cleaner.state.floatbank5_a_level 16775 non-n
ull float64
83 secondary_cleaner.state.floatbank5_b_air 16775 non-n
ull float64
84 secondary_cleaner.state.floatbank5_b_level 16776 non-n
ull float64
85 secondary_cleaner.state.floatbank6_a_air 16757 non-n
ull float64
86 secondary_cleaner.state.floatbank6_a_level 16775 non-n
ull float64
dtypes: float64(86), object(1)
memory usage: 11.2+ MB

```



In [7]:

```
df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5856 entries, 0 to 5855
```

```
Data columns (total 53 columns):
```

#	Column	Non-Null Count	Dty
0	date	5856 non-null	obj
1	primary_cleaner.input.sulfate	5554 non-null	flo
2	primary_cleaner.input.depressant	5572 non-null	flo
3	primary_cleaner.input.feed_size	5856 non-null	flo
4	primary_cleaner.input.xanthate	5690 non-null	flo
5	primary_cleaner.state.floatbank8_a_air	5840 non-null	flo
6	primary_cleaner.state.floatbank8_a_level	5840 non-null	flo
7	primary_cleaner.state.floatbank8_b_air	5840 non-null	flo
8	primary_cleaner.state.floatbank8_b_level	5840 non-null	flo
9	primary_cleaner.state.floatbank8_c_air	5840 non-null	flo
10	primary_cleaner.state.floatbank8_c_level	5840 non-null	flo
11	primary_cleaner.state.floatbank8_d_air	5840 non-null	flo
12	primary_cleaner.state.floatbank8_d_level	5840 non-null	flo
13	rougher.input.feed_ag	5840 non-null	flo
14	rougher.input.feed_pb	5840 non-null	flo
15	rougher.input.feed_rate	5816 non-null	flo
16	rougher.input.feed_size	5834 non-null	flo
17	rougher.input.feed_sol	5789 non-null	flo
18	rougher.input.feed_au	5840 non-null	flo
19	rougher.input.floatbank10_sulfate	5599 non-null	flo
20	rougher.input.floatbank10_xanthate	5733 non-null	flo
21	rougher.input.floatbank11_sulfate	5801 non-null	flo
22	rougher.input.floatbank11_xanthate	5503 non-null	flo
23	rougher.state.floatbank10_a_air	5839 non-null	flo
24	rougher.state.floatbank10_a_level	5840 non-null	flo
25	rougher.state.floatbank10_b_air	5839 non-null	flo
26	rougher.state.floatbank10_b_level	5840 non-null	flo

```

27 rougher.state.floatbank10_c_air          5839 non-null    flo
at64
28 rougher.state.floatbank10_c_level        5840 non-null    flo
at64
29 rougher.state.floatbank10_d_air          5839 non-null    flo
at64
30 rougher.state.floatbank10_d_level        5840 non-null    flo
at64
31 rougher.state.floatbank10_e_air          5839 non-null    flo
at64
32 rougher.state.floatbank10_e_level        5840 non-null    flo
at64
33 rougher.state.floatbank10_f_air          5839 non-null    flo
at64
34 rougher.state.floatbank10_f_level        5840 non-null    flo
at64
35 secondary_cleaner.state.floatbank2_a_air  5836 non-null    flo
at64
36 secondary_cleaner.state.floatbank2_a_level 5840 non-null    flo
at64
37 secondary_cleaner.state.floatbank2_b_air  5833 non-null    flo
at64
38 secondary_cleaner.state.floatbank2_b_level 5840 non-null    flo
at64
39 secondary_cleaner.state.floatbank3_a_air  5822 non-null    flo
at64
40 secondary_cleaner.state.floatbank3_a_level 5840 non-null    flo
at64
41 secondary_cleaner.state.floatbank3_b_air  5840 non-null    flo
at64
42 secondary_cleaner.state.floatbank3_b_level 5840 non-null    flo
at64
43 secondary_cleaner.state.floatbank4_a_air  5840 non-null    flo
at64
44 secondary_cleaner.state.floatbank4_a_level 5840 non-null    flo
at64
45 secondary_cleaner.state.floatbank4_b_air  5840 non-null    flo
at64
46 secondary_cleaner.state.floatbank4_b_level 5840 non-null    flo
at64
47 secondary_cleaner.state.floatbank5_a_air  5840 non-null    flo
at64
48 secondary_cleaner.state.floatbank5_a_level 5840 non-null    flo
at64
49 secondary_cleaner.state.floatbank5_b_air  5840 non-null    flo
at64
50 secondary_cleaner.state.floatbank5_b_level 5840 non-null    flo
at64
51 secondary_cleaner.state.floatbank6_a_air  5840 non-null    flo
at64
52 secondary_cleaner.state.floatbank6_a_level 5840 non-null    flo
at64
dtypes: float64(52), object(1)
memory usage: 2.4+ MB

```

## Notes for preprocessing:

- There are 16860 observations in the train set and 5856 observations in the test set, so the test set is around 26% of the full dataset;
- There are 87 features in the train set and only 53 features in the test set. Some parameters are not available in the test set because they were measured and/or calculated much later. We will analyze this point further in more detail;
- There are 2 target variables: `rougher.output.recovery` and `final.output.recovery`;
- Each feature data type is float except for the `date` column, it should be converted to datetime format;
- There are quite a few missing values to fill.

## Recovery calculation check

Let's use the following formula to check whether `rougher.output.recovery` was calculated correctly in the train set:

$$\text{Recovery} = \frac{C \times (F - T)}{F \times (C - T)} \times 100\%$$

where:

- $C$  — share of gold in the concentrate right after flotation (for finding the rougher concentrate recovery)
- $F$  — share of gold in the feed before flotation (for finding the rougher concentrate recovery)
- $T$  — share of gold in the rougher tails right after flotation (for finding the rougher concentrate recovery)

In [8]:

```
C = df_train['rougher.output.concentrate_au']
F = df_train['rougher.input.feed_au']
T = df_train['rougher.output.tail_au']

df_train['recovery_calc'] = (100 * ((C * (F - T)) / (F * (C - T)))) .round(6)
```

In [9]:

```
df_train[['recovery_calc', 'rougher.output.recovery']][0:5]
```

Out[9]:

	recovery_calc	rougher.output.recovery
0	87.107763	87.107763
1	86.843261	86.843261
2	86.842308	86.842308
3	87.226430	87.226430
4	86.688794	86.688794

The first 5 rows are identical, let's calculate MAE to make sure it holds for all observations.

In [10]:

```
MAE = (df_train['recovery_calc'] - df_train['rougher.output.recovery']).abs().mean()  
MAE
```

Out[10]:

```
2.4482453097445186e-07
```

The MAE value is very close to 0, which means the `rougher.output.recovery` was calculated correctly.

## Missing features analysis

In [11]:

```
missing_col = list(set(df_train.columns)-set(df_test.columns))
sorted(missing_col)
```

Out[11]:

```
['final.output.concentrate_ag',
 'final.output.concentrate_au',
 'final.output.concentrate_pb',
 'final.output.concentrate_sol',
 'final.output.recovery',
 'final.output.tail_ag',
 'final.output.tail_au',
 'final.output.tail_pb',
 'final.output.tail_sol',
 'primary_cleaner.output.concentrate_ag',
 'primary_cleaner.output.concentrate_au',
 'primary_cleaner.output.concentrate_pb',
 'primary_cleaner.output.concentrate_sol',
 'primary_cleaner.output.tail_ag',
 'primary_cleaner.output.tail_au',
 'primary_cleaner.output.tail_pb',
 'primary_cleaner.output.tail_sol',
 'recovery_calc',
 'rougher.calculation.au_pb_ratio',
 'rougher.calculation.floatbank10_sulfate_to_au_feed',
 'rougher.calculation.floatbank11_sulfate_to_au_feed',
 'rougher.calculation.sulfate_to_au_concentrate',
 'rougher.output.concentrate_ag',
 'rougher.output.concentrate_au',
 'rougher.output.concentrate_pb',
 'rougher.output.concentrate_sol',
 'rougher.output.recovery',
 'rougher.output.tail_ag',
 'rougher.output.tail_au',
 'rougher.output.tail_pb',
 'rougher.output.tail_sol',
 'secondary_cleaner.output.tail_ag',
 'secondary_cleaner.output.tail_au',
 'secondary_cleaner.output.tail_pb',
 'secondary_cleaner.output.tail_sol']
```

The missing features are directly connected to the target variables (types **output** and **calculation**). It's only logical that they are not included in the test set as we are developing a model precisely to predict those 2 targets.

In order to be able to use an ML model, we need to remove these extra features from the train set, so that both train and test set have the same shapes.

## Preprocessing

### Data type change

As mentioned above, let's convert the `date` column into the datetime type.



In [12]:

```
df_train['date'] = pd.to_datetime(df_train['date'])
df_test['date'] = pd.to_datetime(df_test['date'])
```

## Test targets

Let's use the full dataset to extract test targets and include them in the test set to be able to compare our predictions to the actual values.

In [13]:

```
df_full['date'] = pd.to_datetime(df_full['date'])
```

In [14]:

```
df_test = df_test.merge(df_full[['date', 'final.output.recovery', 'rougher.output.recovery']],
                        how='left', on='date')
```

## Missing values

In [15]:

```
df_train.isnull().sum().mean()/len(df_train)
```

Out[15]:

0.021974414968187212

In [16]:

```
df_test.isnull().sum().mean()/len(df_test)
```

Out[16]:

0.010394932935916543

In both datasets there are a few missing values in almost each column, but their share is not significant (no more than 2%, on average), we will simply drop them.

In [17]:

```
df_train = df_train.dropna(how='any', axis=0)
df_train.reset_index(drop=True, inplace=True)

df_test = df_test.dropna(how='any', axis=0)
df_test.reset_index(drop=True, inplace=True)
```

In [18]:

```
df_train.isnull().sum().sum()
```

Out[18]:

0

In [19]:

```
df_train.shape
```

Out[19]:

```
(11017, 88)
```

In [20]:

```
df_test.isnull().sum().sum()
```

Out[20]:

```
0
```

In [21]:

```
df_test.shape
```

Out[21]:

```
(5229, 55)
```

All missing values were removed.

## Duplicates

Let's check if any rows are duplicated.

In [22]:

```
df_train.duplicated().sum()
```

Out[22]:

```
0
```

In [23]:

```
df_test.duplicated().sum()
```

Out[23]:

```
0
```

## EDA

### Concentrations of metals

Let's see how the concentrations of metals (Au, Ag, Pb) change depending on the purification stage.

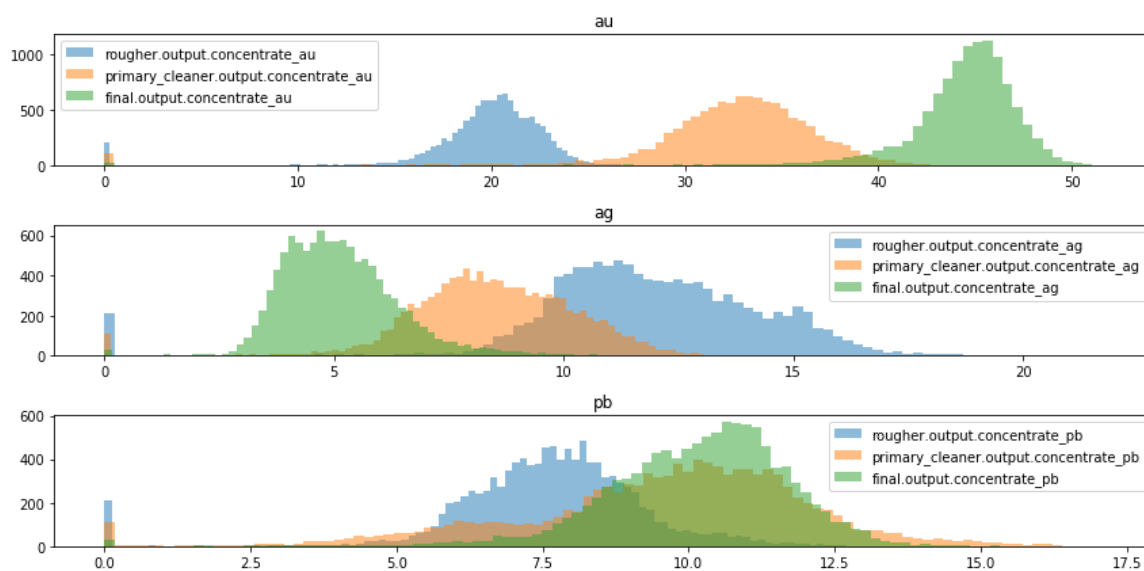
In [24]:

```
metals = ['au', 'ag', 'pb']
stages = ['rougher.output.concentrate', 'primary_cleaner.output.concentrate', 'final.output.concentrate']
plt.figure(figsize=(12,6))

for i,metal in enumerate(metals):
    for stage in stages:
        plt.subplot(3,1,i+1)
        plt.hist(df_train[stage+'_'+metal], bins=100, label=stage+'_'+metal, alpha=.5);
    plt.title(metal)

    plt.legend()

plt.tight_layout()
```



First of all, we see around 2000 outliers for each metal and stage - values with 0 concentration of metals. We will need to remove them as they directly correlate with our targets.

As for the gold concentration ( au ), there is a clear trend towards quality improvement after each stage: the share of gold is higher and higher on average as we proceed with purification (from 20% to more than 45%, on average).

Interestingly, we see the opposite trend for silver concentration ( ag ): the more we purify the feed, the lower gets the share of this metal (from around 11% to less than 5%, on average).

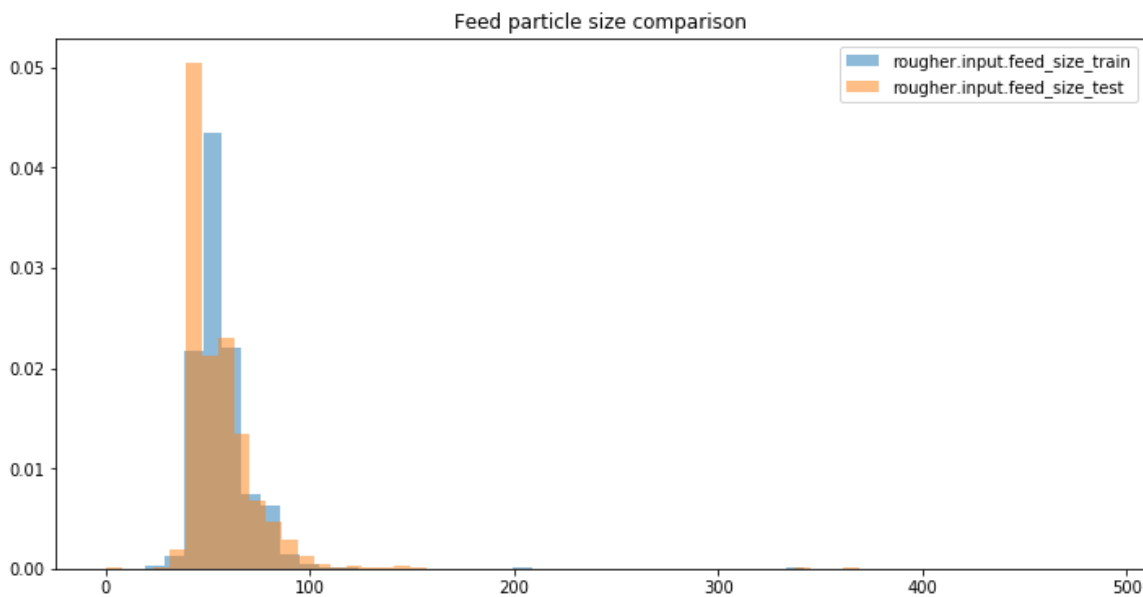
As for the lead concentration ( pb ), we can say that there is probably no need for the second purification stage, as the quality of this metal doesn't improve, on average. However, we see a slight improvement after the first stage (from around 8% to almost 11%, on average).

## Feed particle size comparison

Let's compare the feed particle size distributions in the training set and in the test set. If the distributions vary significantly, the model evaluation will be incorrect.

In [25]:

```
plt.figure(figsize=(12,6))
plt.hist(df_train['rougner.input.feed_size'],bins=50,label='rougner.input.feed_s
ize_train',alpha=.5, density=1)
plt.hist(df_test['rougner.input.feed_size'],bins=50,label='rougner.input.feed_si
ze_test',alpha=.5, density=1)
plt.legend()
plt.title('Feed particle size comparison');
```



The two distributions are very close to each other, which means that we will not have a problem of applying a model trained on the train set to the test set.

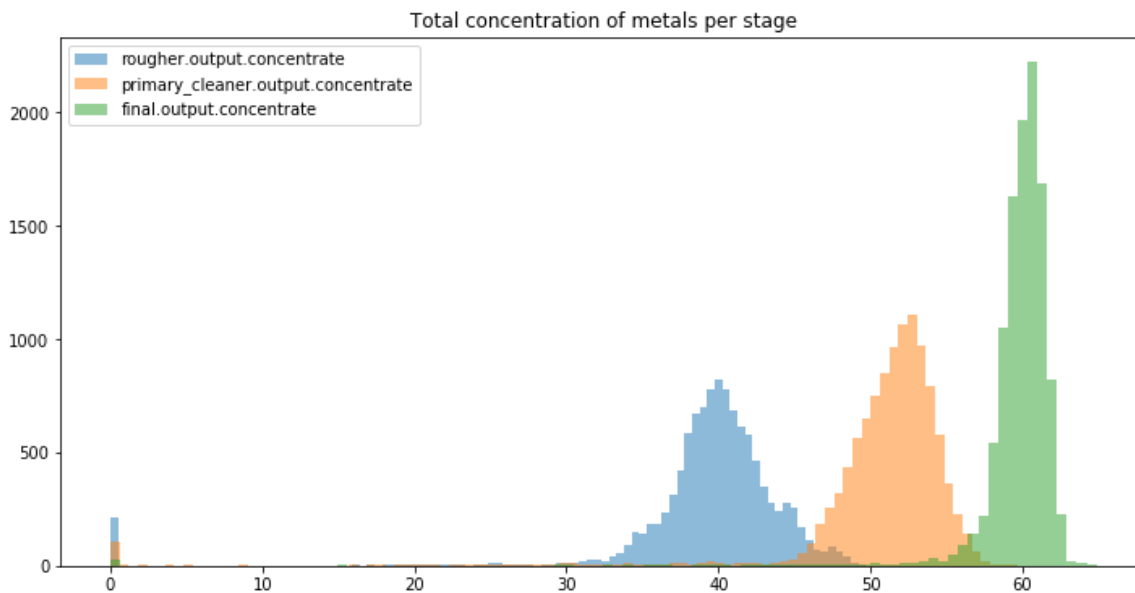
## Outliers

Let's remove the outliers in the target variables. We will first transform any 0 values to NaNs and then drop them from the data sets because their amount is insignificant.

In [26]:

```
plt.figure(figsize=(12,6))

for stage in stages:
    df_train[stage+'_all_metals'] = 0
    for metal in metals:
        df_train[stage+'_all_metals'] += df_train[stage+'_'+metal]
    plt.hist(df_train[stage+'_all_metals'], bins=100, label=stage, alpha=.5)
    plt.title('Total concentration of metals per stage')
    plt.legend();
```



We can see that the total concentration of metals per stage does get better, on average.

Just as before we see multiple outliers around 0, let's remove them.

In [27]:

```
for i,metal in enumerate(metals):
    for stage in stages:
        df_train[df_train[stage+'_'+metal]<0.01] = np.nan
df_train = df_train.dropna(how='any', axis=0)
```

In [28]:

```
df_train.isnull().sum().sum()
```

Out[28]:

0

In [29]:

```
df_train.shape
```

Out[29]:

```
(10676, 91)
```

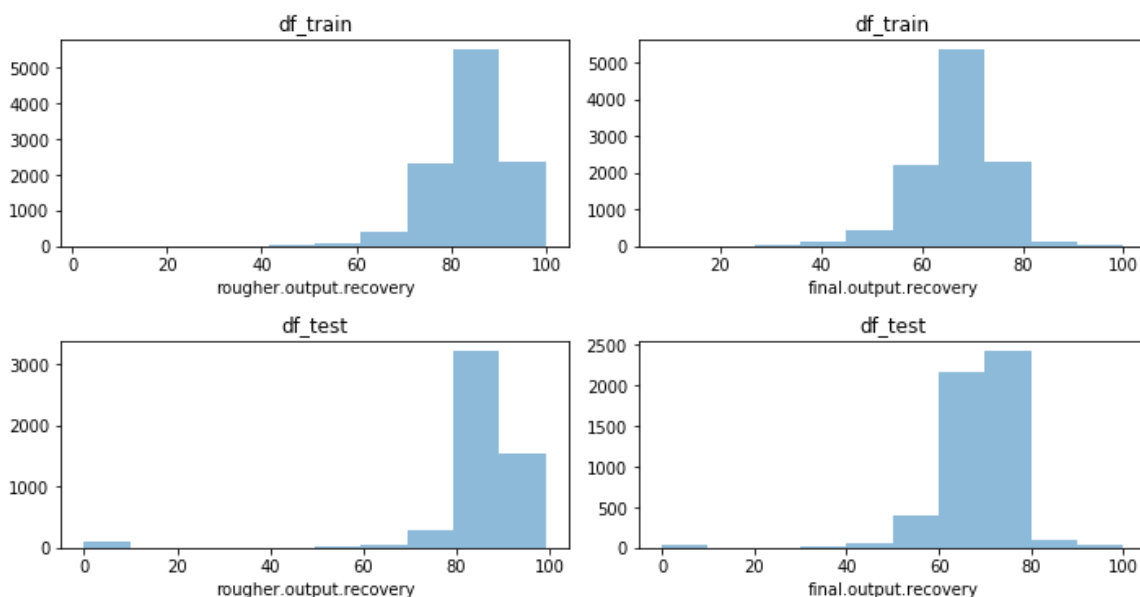
## Target analysis

Finally, let's analyze our targets.

In [30]:

```
targets = ['rougner.output.recovery', 'final.output.recovery']
dfs = [df_train, df_test]
df_train.name = 'df_train'
df_test.name = 'df_test'

plt.figure(figsize=(10,10))
c=1
for i, df in enumerate(dfs):
    for target in targets:
        plt.subplot(4,2,c)
        plt.hist(df[target], alpha=.5)
        plt.title(df.name)
        plt.xlabel(target)
        c = c + 1
plt.tight_layout()
```



Train distribution looks ok, both targets in the test set have outliers around 0, let's remove them.

In [31]:

```
for target in targets:
    df_test[df_test[target]<0.01] = np.nan

df_test = df_test.dropna(how='any', axis=0)
```

In [32]:

```
df_test.isnull().sum().sum()
```

Out[32]:

0

In [33]:

```
df_test.shape
```

Out[33]:

(5105, 55)

## Common columns

In the end, we will get rid of extra features in the train set.

In [34]:

```
common_columns = list(set(df_train.columns).intersection(set(df_test.columns)))
df_train_filtered = df_train[common_columns]
df_train_filtered[['rougher.output.recovery', 'final.output.recovery']] = df_train[['rougher.output.recovery', 'final.output.recovery']]
```

## Standard scaling

Let's scale the features before modeling to be able to compare their coefficients in the later sections.

In [35]:

```
X_train = df_train_filtered.drop(['rougher.output.recovery', 'final.output.recovery', 'date'], axis=1)
X_test = df_test.drop(['rougher.output.recovery', 'final.output.recovery', 'date'], axis=1)
y_train = df_train_filtered[['rougher.output.recovery', 'final.output.recovery']].values
y_test = df_test[['rougher.output.recovery', 'final.output.recovery']].values

sc = ss()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## Evaluation metric

Let's write a function to calculate the final sMAPE value.

**sMAPE** is a symmetric Mean Absolute Percentage Error.

It is similar to MAE, but is expressed in relative values instead of absolute ones. It equally takes into account the scale of both the target and the prediction.

Here's how *sMAPE* is calculated:

$$\text{sMAPE} = \frac{1}{N} \sum_{i=1}^N \frac{|y_i - \hat{y}_i|}{(|y_i| + |\hat{y}_i|) / 2} \times 100\%$$

Denotation:

- Value of target for the observation with the  $i$  index in the sample used to measure quality.
- Value of prediction for the observation with the  $i$  index, for example, in the test sample.
- Number of observations in the sample.
- Summation over all observations of the sample ( $i$  takes values from 1 to  $N$ ).

We need to predict two values:

1. rougher concentrate recovery `rougher.output.recovery`
2. final concentrate recovery `final.output.recovery`

The final metric includes the two values:

$$\text{Final sMAPE} = 25\% \times \text{sMAPE(rougher)} + 75\% \times \text{sMAPE(final)}$$

In [36]:

```
def smape(y_true, y_pred):
    return (np.abs(y_true-y_pred)/((np.abs(y_true) + np.abs(y_pred))/2)).mean()
```

In [37]:

```
def smape_final(y_true,y_pred):
    smape_rougher = smape(y_true[:,0], y_pred[:,0])
    smape_final = smape(y_true[:,1], y_pred[:,1])
    return 0.25*smape_rougher + 0.75*smape_final
```

## Model selection

Now let's train our models on the train set and select the best model using the cross-validation technique.



In [38]:

```
LR = LinearRegression()
DT = DecisionTreeRegressor(random_state=12345)
RF = RandomForestRegressor(random_state=12345)
Lasso = linear_model.Lasso()
Ridge = linear_model.Ridge()
base_model = DummyRegressor(strategy='mean')
```

In [39]:

```
def crossval(model, X_train, y_train, cv):
    smape_score = make_scorer(smape_final)
    scores = cross_val_score(model, X_train, y_train, cv=cv, scoring=smape_score)
    return scores.mean()
```

In [40]:

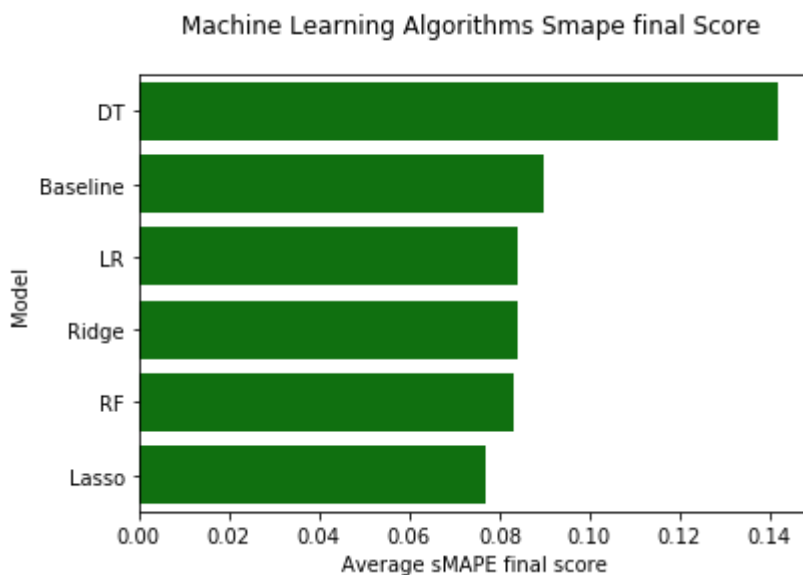
```
smape_final_Lasso = crossval(Lasso, X_train, y_train, 5)
smape_final_Ridge = crossval(Ridge, X_train, y_train, 5)
smape_final_LR = crossval(LR, X_train, y_train, 5)
smape_final_DT = crossval(DT, X_train, y_train, 5)
smape_final_RF = crossval(RF, X_train, y_train, 5)
smape_final_base = crossval(base_model, X_train, y_train, 5)
```

In [41]:

```
models = pd.DataFrame({
    'Model': ['Baseline', 'LR', 'DT', 'RF', 'Lasso', 'Ridge'],
    'Score': [smape_final_base, smape_final_LR, smape_final_DT, smape_final_RF,
smape_final_Lasso, smape_final_Ridge]})
sorted_by_score = models.sort_values(by='Score', ascending=False)
```

In [42]:

```
sns.barplot(x='Score', y = 'Model', data = sorted_by_score, color = 'g')
plt.title('Machine Learning Algorithms Smape final Score \n')
plt.xlabel('Average sMAPE final score')
plt.ylabel('Model');
```



Lasso regression model shows the best average score. Let's try to tune its hyperparameters.

## Hyperparameter tuning

In [43]:

```
d = []
for alpha in np.arange(0.1,1,0.1):
    for tol in [0.01, 0.05, 0.001]:
        for max_iter in (100, 500, 1000):
            Lasso = linear_model.Lasso(alpha=alpha, tol=tol, max_iter=max_iter)
            smape_final_Lasso = crossval(Lasso, X_train, y_train, 5)
            d.append(
                {
                    'alpha': alpha,
                    'tol': tol,
                    'max_iter': max_iter,
                    'smape_final_Lasso': smape_final_Lasso
                }
            )

best_param = pd.DataFrame(d).nlargest(1, ['smape_final_Lasso'], keep='first')
smape_final_RF = best_param['smape_final_Lasso'].values
best_param
```

Out[43]:

	alpha	tol	max_iter	smape_final_Lasso
3	0.1	0.05	100	0.08001

## Test the model

In [44]:

```
Lasso = linear_model.Lasso(alpha=0.1, tol=0.001, max_iter=100)
Lasso.fit(X_train, y_train)
y_pred = Lasso.predict(X_test)
smape_final(y_test, y_pred)
```

Out[44]:

2.0

The test score is pretty bad even after hyper parameters tuning. The model seems to be overfitted to the train set. Let's try feature selection to reduce overfitting.

## Feature importance

One of the goals of this report is to eliminate unprofitable features. We can do that by comparing feature importances.

In [45]:

```
LR.fit(X_train, y_train)
coeff_df = pd.DataFrame()
coeff_df['Feature'] = df_train_filtered.drop(['rougher.output.recovery', 'final.o
utput.recovery', 'date'], axis=1).columns.values
coeff_df["Correlation"] = pd.Series(LR.coef_[0])

coeff_df.sort_values(by='Correlation', ascending=False)
```

Out[45]:

	Feature	Correlation
50	rougher.state.floatbank10_f_air	2.873959
41	rougher.input.feed_ag	1.969526
29	rougher.state.floatbank10_b_level	1.882530
9	rougher.input.feed_sol	1.829609
30	rougher.state.floatbank10_a_level	1.695842
20	primary_cleaner.input.sulfate	1.635028
43	secondary_cleaner.state.floatbank5_a_air	1.521347
35	rougher.input.floatbank11_sulfate	1.245344
38	primary_cleaner.state.floatbank8_b_air	1.223162
17	secondary_cleaner.state.floatbank3_a_air	1.167750
42	rougher.input.floatbank10_xanthate	1.153397
33	secondary_cleaner.state.floatbank2_a_air	1.106496
45	secondary_cleaner.state.floatbank5_b_level	1.069566
7	primary_cleaner.state.floatbank8_c_level	0.695207
36	secondary_cleaner.state.floatbank3_b_level	0.582116
3	secondary_cleaner.state.floatbank4_b_level	0.532246
6	rougher.input.floatbank11_xanthate	0.347046
13	primary_cleaner.input.feed_size	0.299679
46	rougher.input.feed_pb	0.237812
34	secondary_cleaner.state.floatbank2_a_level	0.175052
37	secondary_cleaner.state.floatbank4_a_level	0.161707
4	secondary_cleaner.state.floatbank4_b_air	0.125723
16	rougher.state.floatbank10_b_air	0.060345
51	rougher.input.feed_au	0.049011
21	primary_cleaner.state.floatbank8_d_level	0.047284
28	rougher.input.feed_size	-0.015142
18	rougher.state.floatbank10_f_level	-0.116173
47	primary_cleaner.input.xanthate	-0.136632
39	primary_cleaner.state.floatbank8_a_air	-0.140311
5	secondary_cleaner.state.floatbank5_b_air	-0.175401
27	primary_cleaner.input.depressant	-0.200300
15	primary_cleaner.state.floatbank8_b_level	-0.247245
40	rougher.state.floatbank10_d_level	-0.258105
0	primary_cleaner.state.floatbank8_a_level	-0.261074
10	secondary_cleaner.state.floatbank5_a_level	-0.357582
11	primary_cleaner.state.floatbank8_c_air	-0.359393

	Feature	Correlation
26	secondary_cleaner.state.floatbank4_a_air	-0.384321
23	secondary_cleaner.state.floatbank2_b_level	-0.545032
32	rougher.state.floatbank10_d_air	-0.593025
22	rougher.state.floatbank10_e_level	-0.630099
12	secondary_cleaner.state.floatbank2_b_air	-0.705941
24	secondary_cleaner.state.floatbank6_a_air	-0.742980
44	secondary_cleaner.state.floatbank6_a_level	-0.753929
31	rougher.state.floatbank10_c_air	-0.773621
19	primary_cleaner.state.floatbank8_d_air	-0.988076
49	secondary_cleaner.state.floatbank3_a_level	-1.069232
1	secondary_cleaner.state.floatbank3_b_air	-1.101443
25	rougher.state.floatbank10_c_level	-1.157193
48	rougher.state.floatbank10_a_air	-1.270705
14	rougher.input.feed_rate	-1.903286
8	rougher.state.floatbank10_e_air	-1.986217
2	rougher.input.floatbank10_sulfate	-2.558315

From the above table we can see that `rougher.state.floatbank10_e_air` and `rougher.state.floatbank10_f_air` have the highest impact on the target variables. The features with scores close to 0 have almost no impact on the targets and can be eliminated.

In [46]:

```
profitable_params = coeff_df.loc[abs(coeff_df["Correlation"]) > 0.1 , 'Feature']
```

In [47]:

```
df_train_params = df_train_filtered[profitable_params]
df_test_params = df_test[profitable_params]

df_train_params[['rougher.output.recovery', 'final.output.recovery']] = df_train_filtered[['rougher.output.recovery', 'final.output.recovery']]
df_test_params[['rougher.output.recovery', 'final.output.recovery']] = df_test[['rougher.output.recovery', 'final.output.recovery']]
```

In [48]:

```
X_train = df_train_params.drop(['rougner.output.recovery', 'final.output.recover
y'], axis=1)
X_test = df_test_params.drop(['rougner.output.recovery', 'final.output.recover
y'], axis=1)
y_train = df_train_params[['rougner.output.recovery', 'final.output.recovery']].
values
y_test = df_test_params[['rougner.output.recovery', 'final.output.recovery']].va
lues

sc = ss()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

In [49]:

```
Lasso = linear_model.Lasso(alpha=0.1, tol=0.001, max_iter=100)
Lasso.fit(X_train, y_train)
y_pred = Lasso.predict(X_test)
smape_final_score = smape_final(y_test, y_pred)
smape_final_score
```

Out[49]:

0.06794437817892128

## Sanity check

Let's calculate the test baseline score to perform this check.

In [50]:

```
base_model = DummyRegressor(strategy='mean')
base_model.fit(X_train, y_train)
y_pred = base_model.predict(X_test)
smape_final_base = smape_final(y_test, y_pred)
smape_final_base
```

Out[50]:

0.07637178944673645

In [51]:

```
print(round((smape_final_base - smape_final_score)/smape_final_base,2)*100,'%')
```

11.0 %

Lasso model cross-validation score is lower than the baseline score by 11%.

## Conclusion

In this project we have **developed a machine learning model for Zyfra, efficiency solutions for heavy industry company, that analyzes data on extraction and purification of gold from an ore and predicts the amount of recovered gold. The model helps to eliminate unprofitable parameters.**

First of all, we have familiarized ourselves with the data by performing the descriptive statistics. We found missing values, wrong `date` column data type.

Next, we checked the recovery calculation in the dataset. It turned out to be performed correctly as the MAE value between the `rougher.output.recovery` column and our calculations was very close to 0.

We found that some features are missing from the test set as they are directly connected to the target variables (types output and calculation).

In the preprocessing step we have converted `date` column data type to `datetime`, filled missing tenure values with the mean per each column and checked for duplicated values. In order to be able to use an ML model, we removed those extra features from the train set, so that both train and test set have the same shapes.

In the following section we have performed an **exploratory data analysis** and reached the following conclusions:

- As for the gold concentration (au), there is a clear trend towards quality improvement after each stage: the share of gold is higher and higher on average as we proceed with purification (from 20% to more than 45%, on average).
- Interestingly, we see the opposite trend for silver concentration (ag): the more we purify the feed, the lower gets the share of this metal (from around 11% to less than 5%, on average).
- As for the lead concentration (pb), we can say that there is probably no need for the second purification stage, as the quality of this metal doesn't improve, on average. However, we see a slight improvement after the first stage (from around 8% to almost 11%, on average).
- The train and test distributions of feed particle sizes are very close to each other, which means that we will not have a problem of applying a model trained on the train set to the test set.
- We found some outliers in the target variables and removed them.

In the next step we developed and tested, using cross-validation, several ML algorithms and tuned the best model's hyperparameters. Lasso regression model showed the lowest score.

The selected model appeared to be overfitting to the train set, so we decided to implement feature selection method to reduce the overfitting. We have measured feature importances using the coefficients from the Linear Regression model and identified the most profitable parameters.

Then we have tested our model on the test set. We have reached 6.8% error rate on the test set.

Finally, we have checked our model for sanity by comparing the final score to the baseline score. The final score of our model is lower (by 11%) than the baseline score.